



DeepLearning.AI



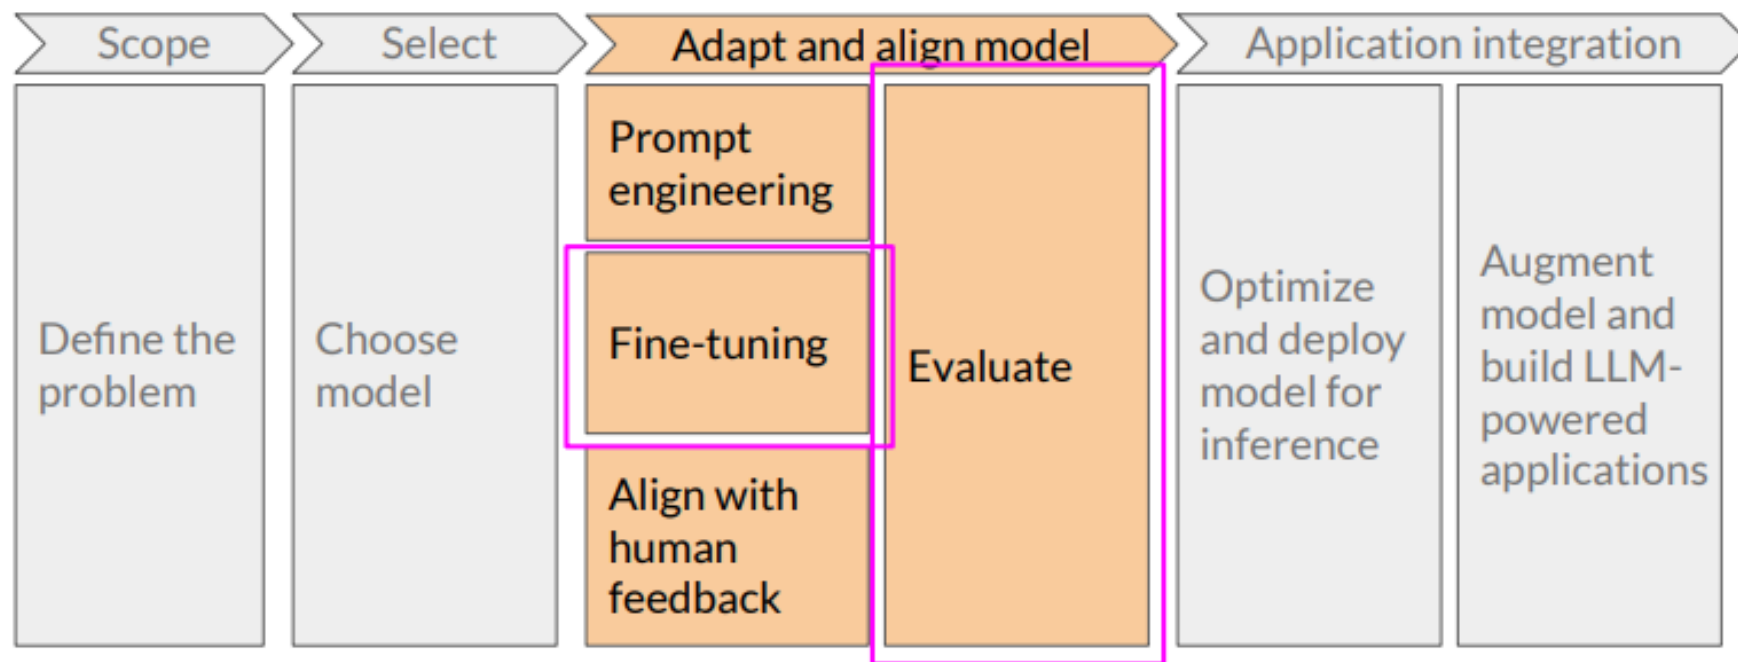
Generative AI and large-language models (LLMs)

**FINE-TUNING, INSTRUCTION
PROMPTS, AND PARAMETER
EFFICIENT FINE-TUNING**

Fine-tuning with instruction prompts




GenAI project lifecycle




Fine-tuning an LLM with instruction prompts

In-context learning (ICL) - zero shot inference



In-context learning (ICL) - one/few shot inference



One-shot or Few-shot Inference

Limitations of in-context learning



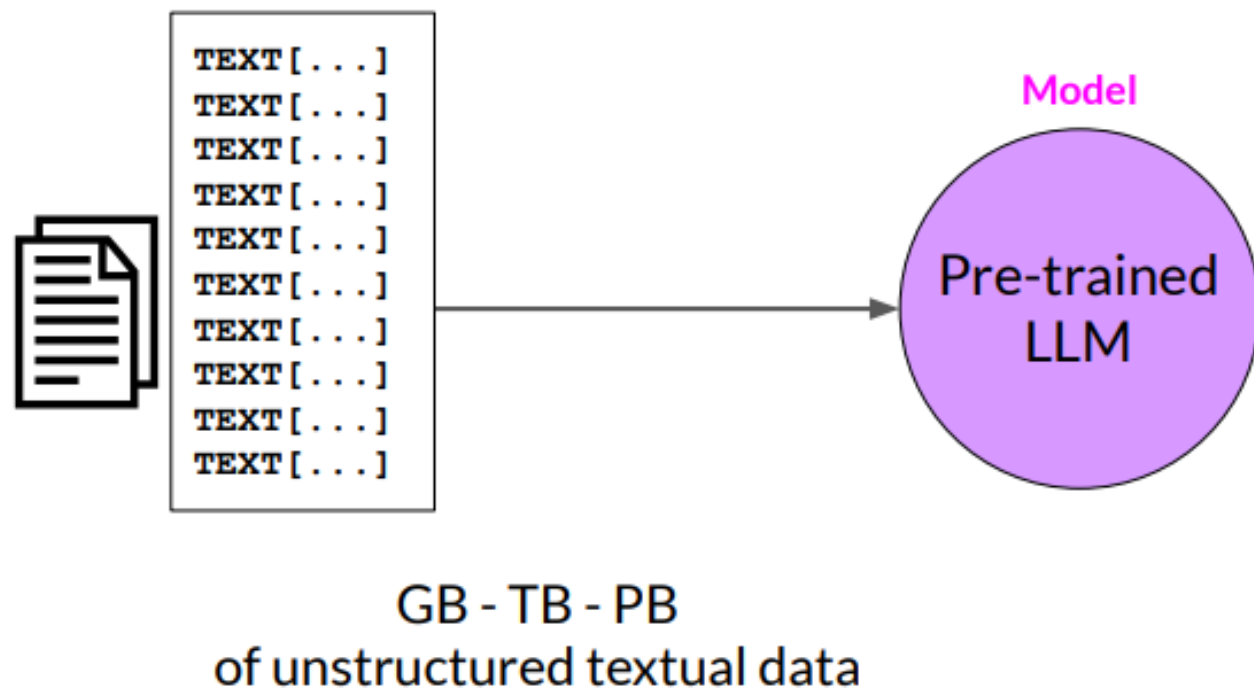
Even with
multiple
examples

- In-context learning may not work for smaller models **LLM**
- Examples take up space in the context window

Instead, try **fine-tuning** the model

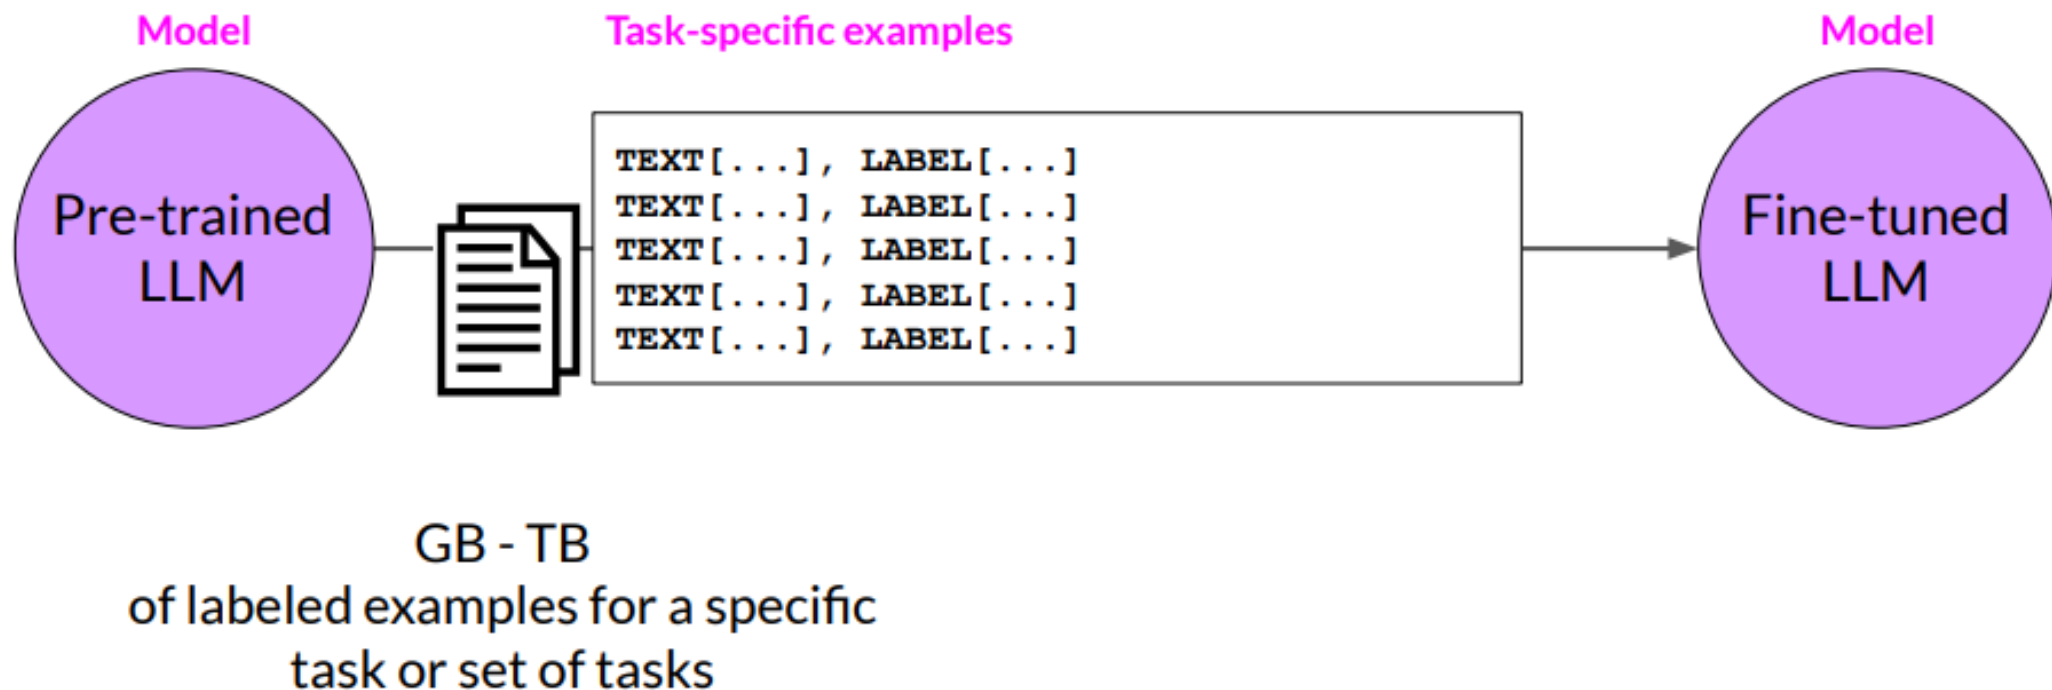
LLM fine-tuning at a high level

LLM pre-training



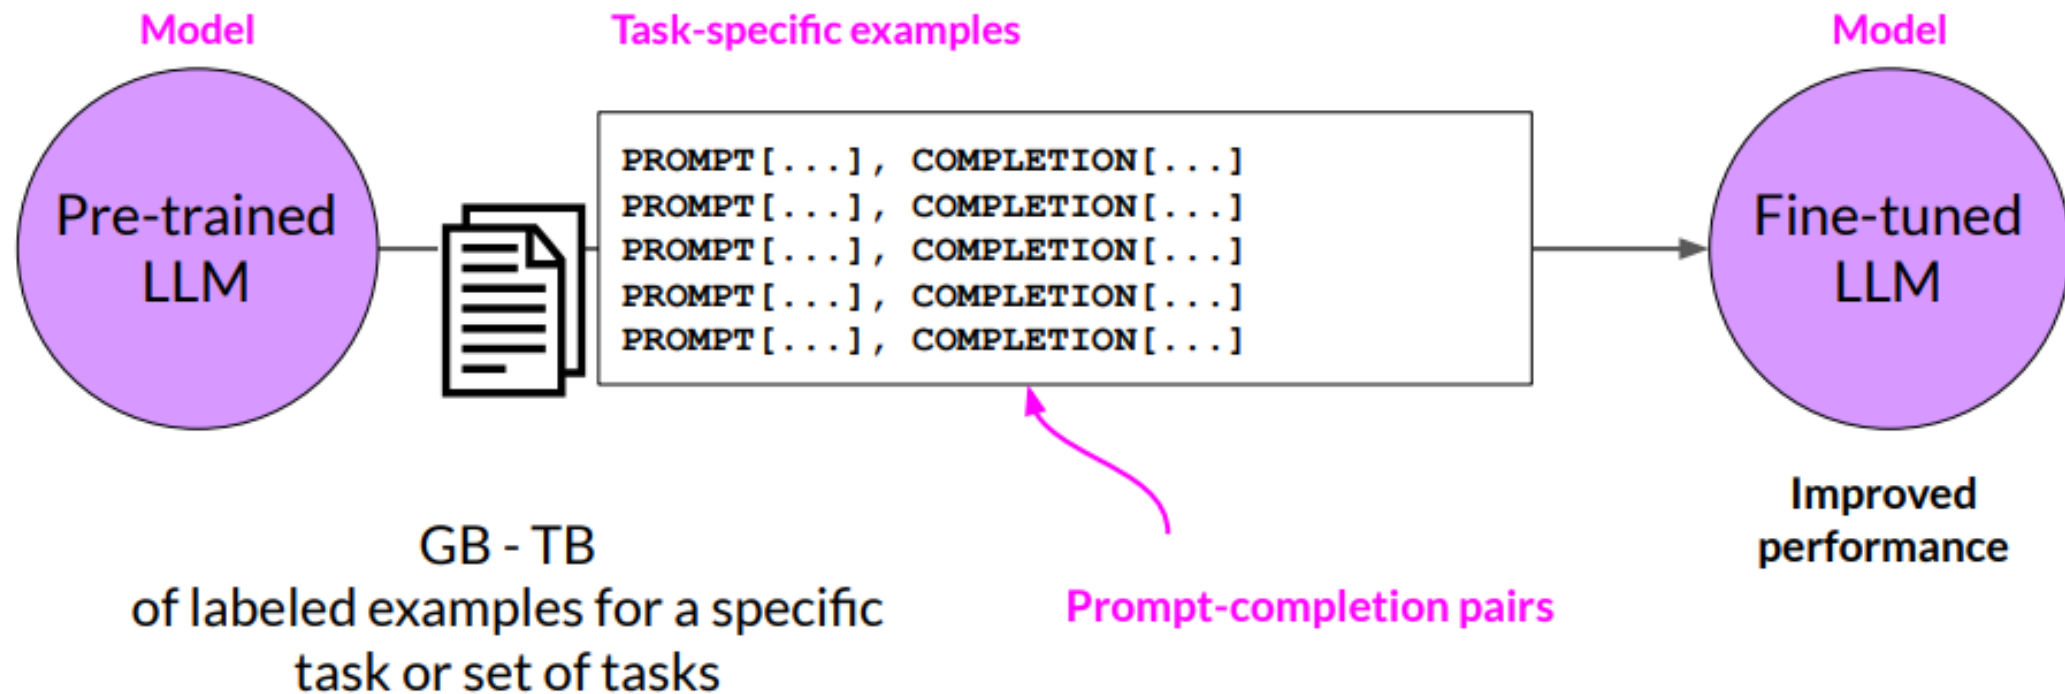
LLM fine-tuning at a high level

LLM fine-tuning



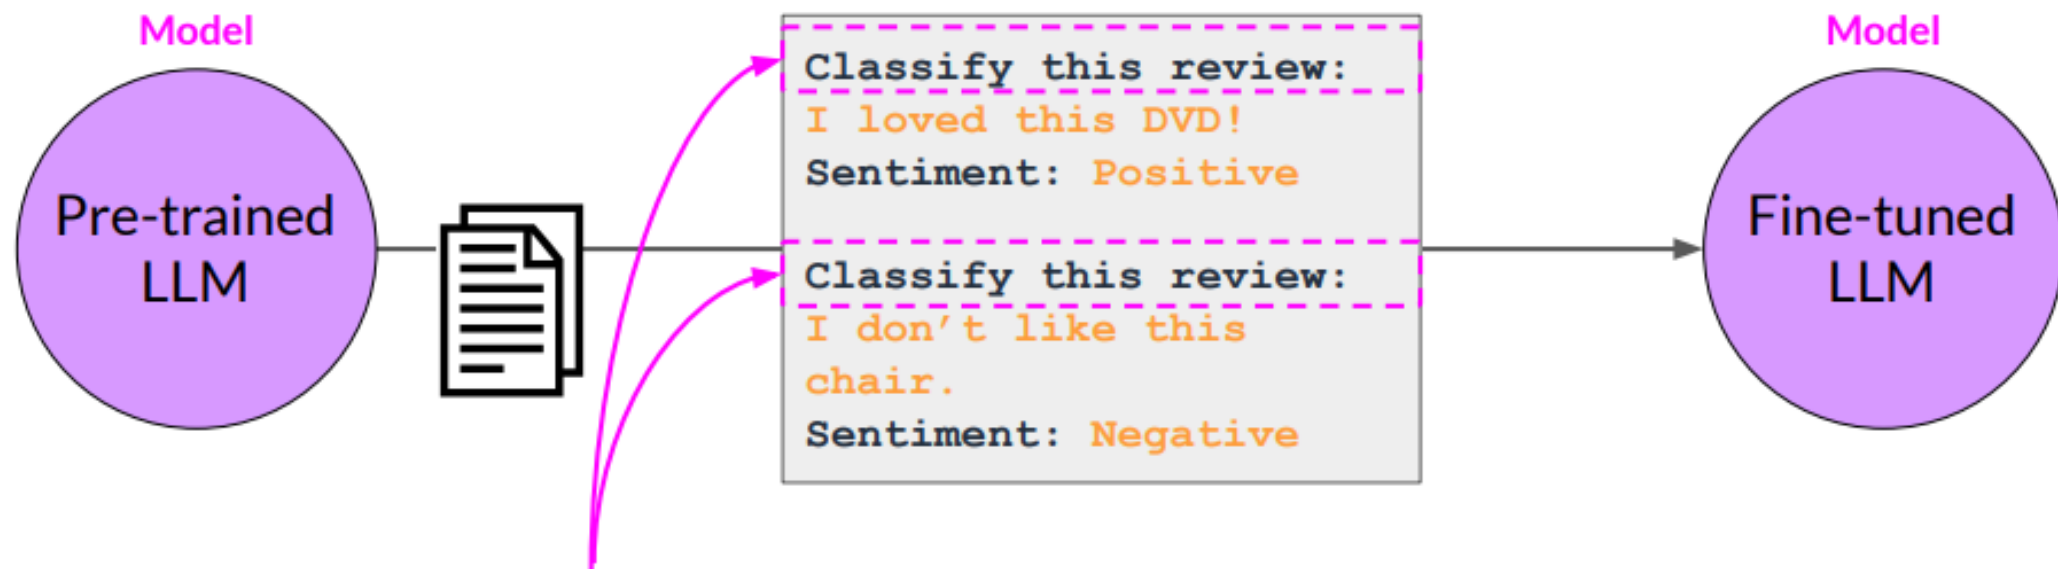
LLM fine-tuning at a high level

LLM fine-tuning



Using prompts to fine-tune LLMs with instruction

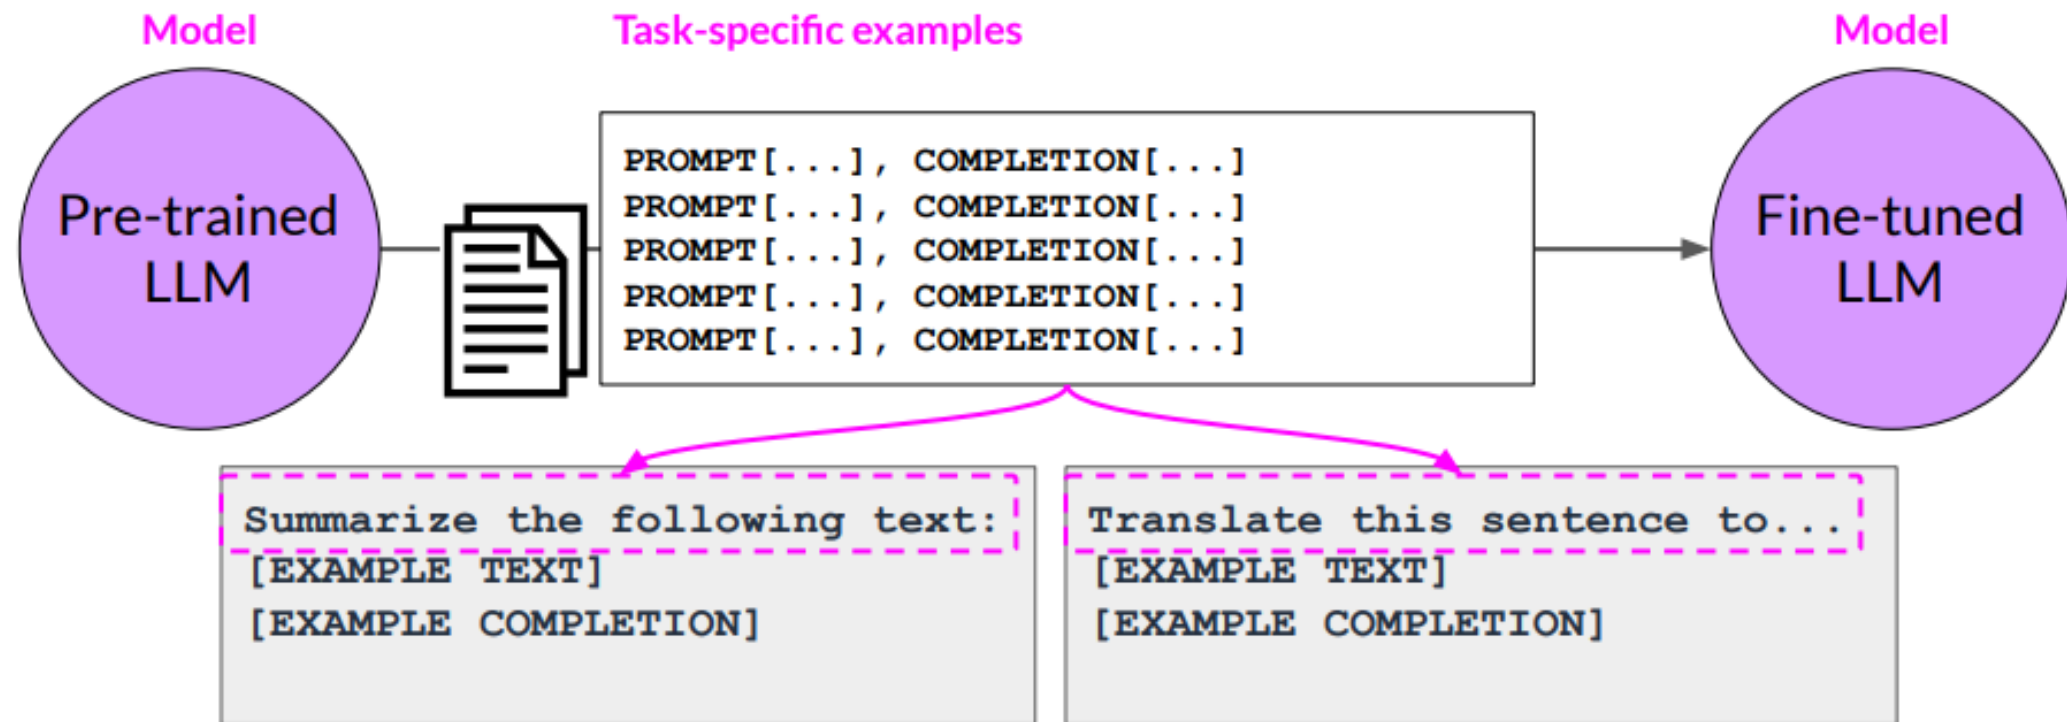
LLM fine-tuning



Each prompt/completion pair includes a specific "instruction" to the LLM

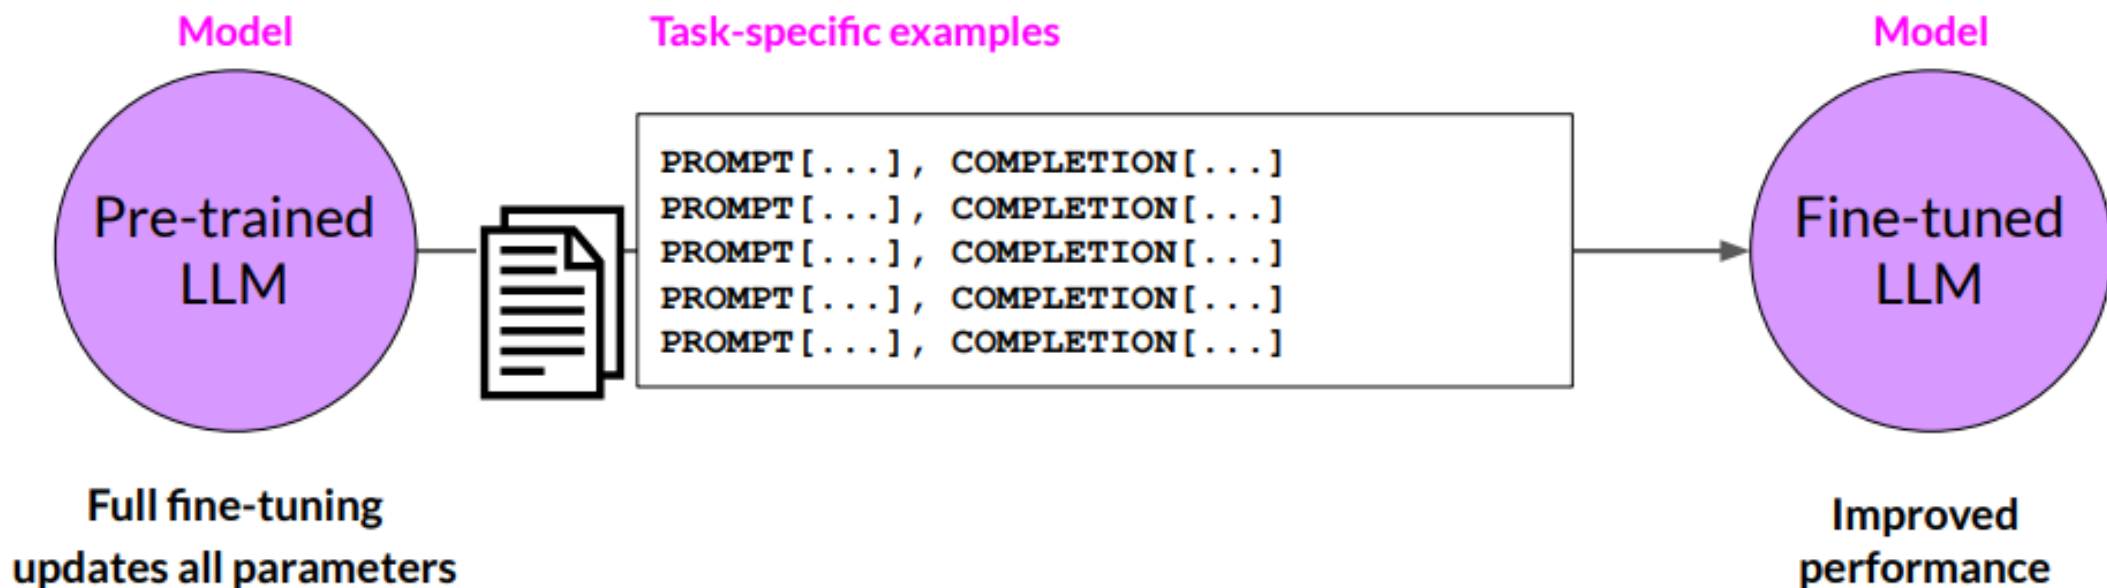
Using prompts to fine-tune LLMs with instruction

LLM fine-tuning



Using prompts to fine-tune LLMs with instruction

LLM fine-tuning



Sample prompt instruction templates

Classification / sentiment analysis

```
jinja: "Given the following review:\n{{review_body}}\n\npredict the associated rating\n\ from the following choices (1 being lowest and 5 being highest)\n- {{ answer_choices\n\ | join('\n- ') }} \n|||\n\n{{answer_choices[star_rating-1]}}"
```

Text generation

```
jinja: Generate a {{star_rating}}-star review (1 being lowest and 5 being highest)\nabout this product {{product_title}}. ||| {{review_body}}
```

Text summarization

```
jinja: "Give a short sentence describing the following product review:\n{{review_body}}\n\n|||\n\n{{review_headline}}"
```

Source: https://github.com/bigscience-workshop/promptsources/blob/main/promptsources/templates/amazon_polarity/templates.yaml

LLM fine-tuning process

LLM fine-tuning

Prepared instruction dataset



Training splits

```
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]
```

Training

```
PROMPT [...], COMPLETION [...]  
...
```

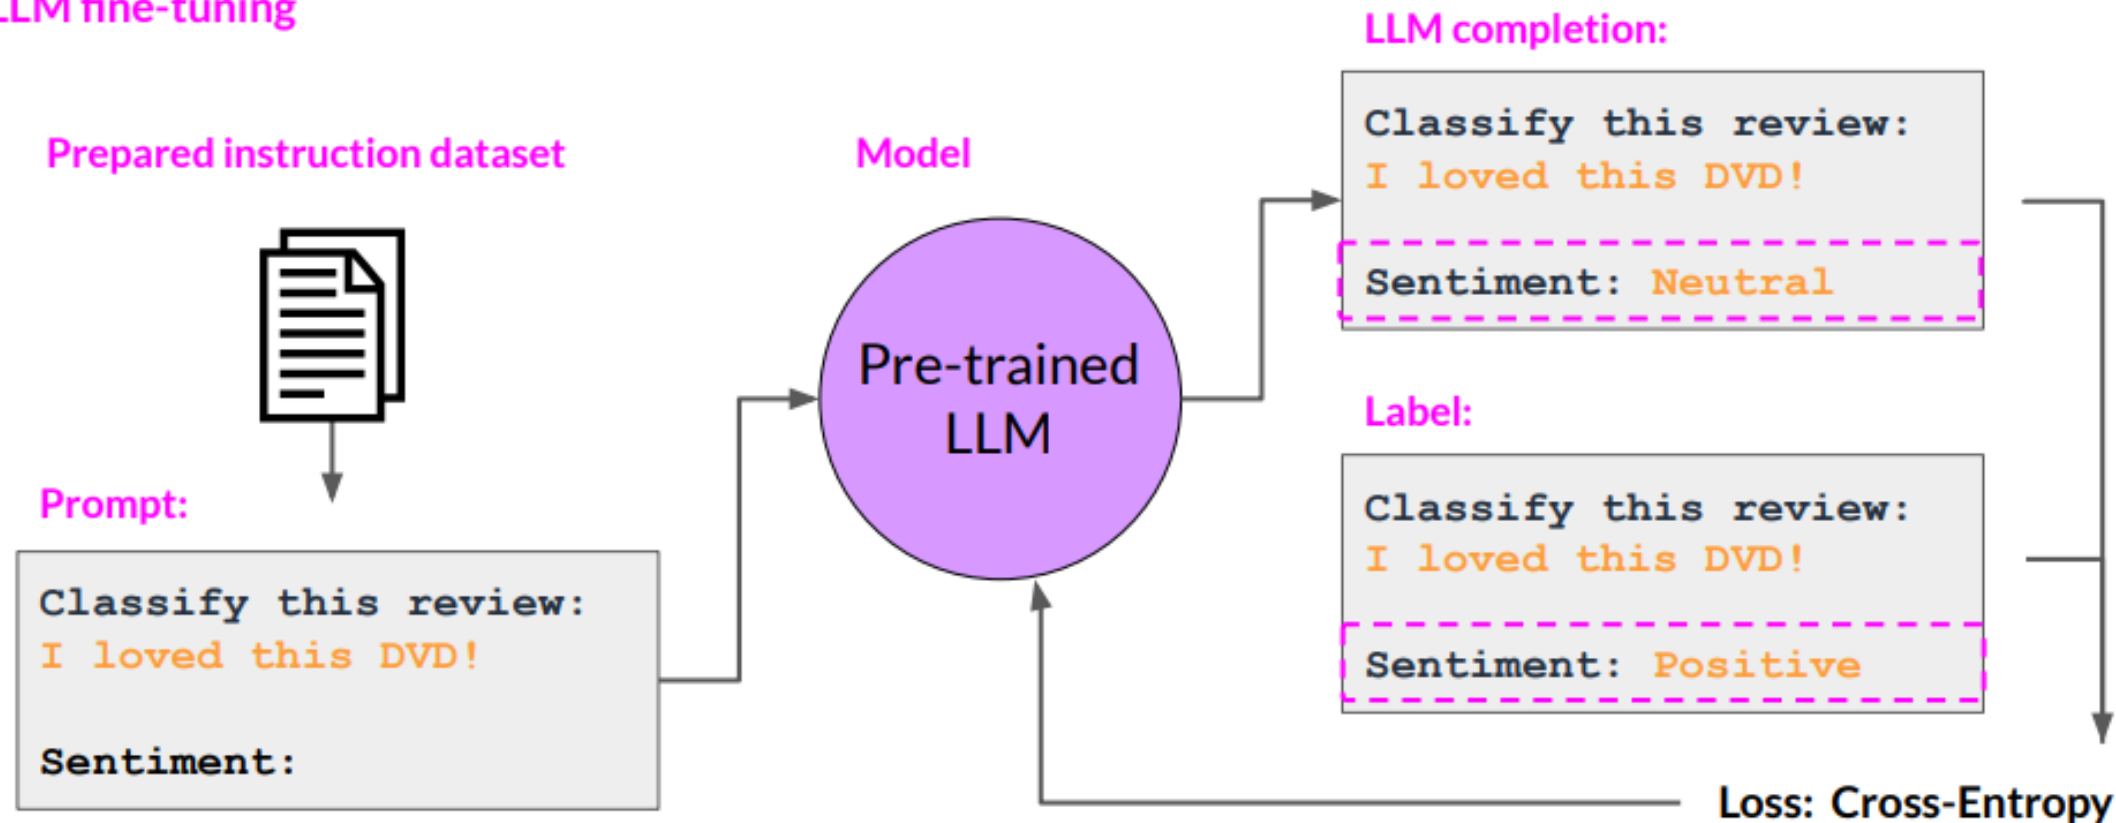
Validation

```
PROMPT [...], COMPLETION [...]  
...
```

Test

LLM fine-tuning process

LLM fine-tuning



LLM fine-tuning process

LLM fine-tuning

Prepared instruction dataset



Training splits

```
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]
```

Training

```
PROMPT [...], COMPLETION [...]  
...
```

Validation

validation_accuracy

```
PROMPT [...], COMPLETION [...]  
...
```

Test

LLM fine-tuning process

LLM fine-tuning

Prepared instruction dataset



Training splits

```
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]  
PROMPT [...], COMPLETION [...]
```

Training

```
PROMPT [...], COMPLETION [...]  
...
```

Validation

```
PROMPT [...], COMPLETION [...]  
...
```

Test

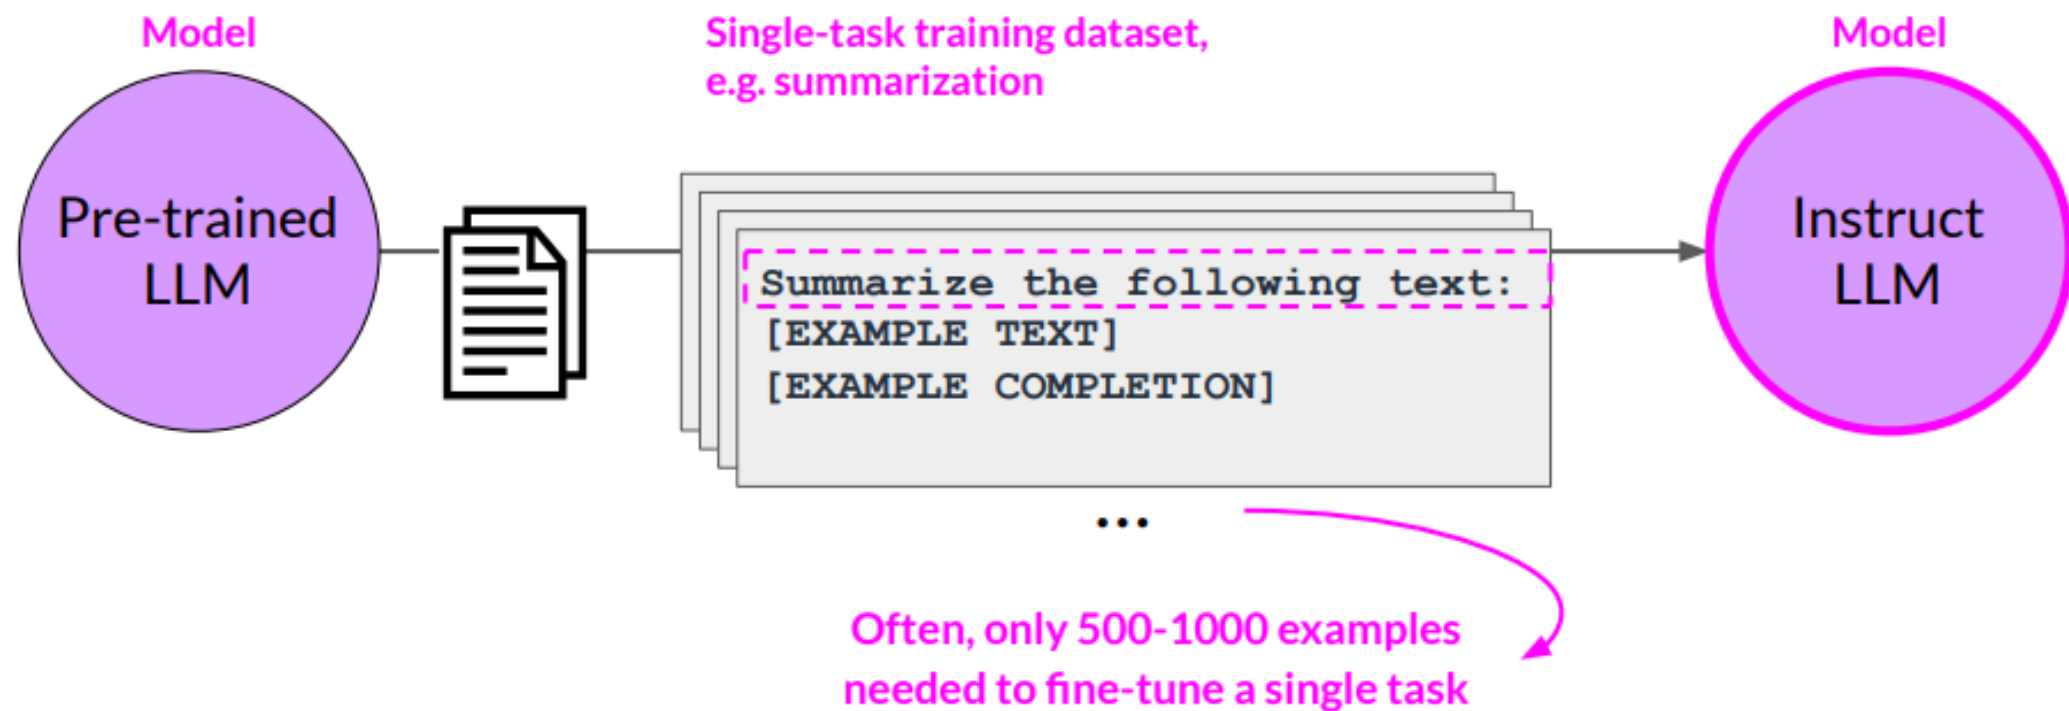
test_accuracy

LLM fine-tuning process



Fine-tuning on a single task

Fine-tuning on a single task



Catastrophic forgetting

- Fine-tuning can significantly increase the performance of a model on a specific task...

Before fine-tuning

Prompt

```
Classify this review:  
I loved this DVD!  
Sentiment:
```

Model



Completion

```
Classify this review:  
I loved this DVD!  
Sentiment: eived a  
very nice book review
```

Catastrophic forgetting

- Fine-tuning can significantly increase the performance of a model on a specific task...

After fine-tuning

Prompt

```
Classify this review:  
I loved this DVD!  
Sentiment:
```

Model



Completion

```
Classify this review:  
I loved this DVD!  
Sentiment: POSITIVE
```

Catastrophic forgetting

- ...but can lead to reduction in ability on other tasks

Before fine-tuning

Prompt

What is the name of
the cat?
Charlie the cat roamed
the garden at night.

Model



Completion

What is the name of
the cat?
Charlie the cat roamed
the garden at night.
Charlie

Catastrophic forgetting

- ...but can lead to reduction in ability on other tasks

After fine-tuning

Prompt

What is the name of
the cat?
Charlie the cat roamed
the garden at night.

Model



Completion

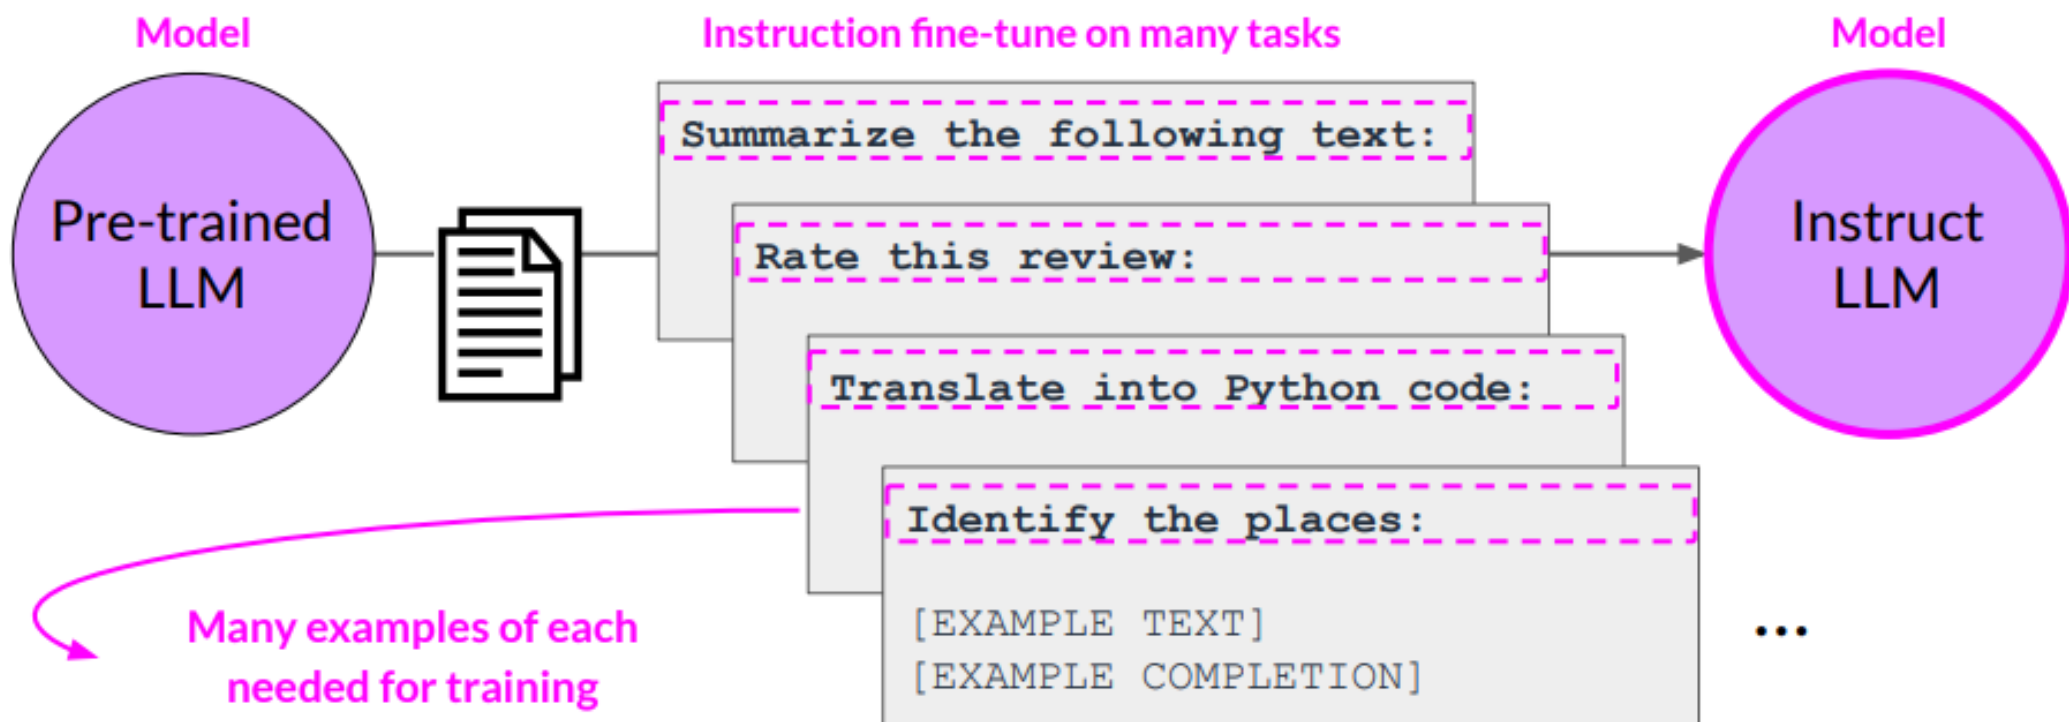
What is the name of
the cat?
Charlie the cat roamed
the garden at night.
**The garden was
positive.**

How to avoid catastrophic forgetting

- First note that you might not have to!
- Fine-tune on **multiple tasks** at the same time
- Consider **Parameter Efficient Fine-tuning** (PEFT)


Multi-task, instruction fine-tuning

Multi-task, instruction fine-tuning



Instruction fine-tuning with FLAN

- FLAN models refer to a specific set of instructions used to perform instruction fine-tuning



FLAN-T5: Fine-tuned version of pre-trained T5 model

- FLAN-T5 is a great, general purpose, instruct model

T0-SF

- Commonsense Reasoning,
- Question Generation,
- Closed-book QA,
- Adversarial QA,
- Extractive QA

...

55 Datasets
14 Categories
193 Tasks

Muffin

- Natural language inference,
- Code instruction gen,
- Code repair
- ~~Dialog context generation,~~
- **Summarization (SAMSum)**

...

69 Datasets
27 Categories
80 Tasks

CoT (reasoning)

- Arithmetic reasoning,
- Commonsense reasoning
- Explanation generation,
- Sentence composition,
- Implicit reasoning,

...

9 Datasets
1 Category
9 Tasks

Natural Instructions

- Cause effect classification,
- Commonsense reasoning,
- Named Entity Recognition,
- Toxic Language Detection,
- Question answering




...

372 Datasets
108 Categories
1554 Tasks

Source: Chung et al. 2022, "Scaling Instruction-Finetuned Language Models"

SAMSum: A dialogue dataset

Sample prompt training dataset (**samsum**) to fine-tune FLAN-T5 from pretrained T5

 Datasets: samsum		Tasks:  Summarization	Languages:  English
dialogue (string)	summary (string)		
"Amanda: I baked cookies. Do you want some? Jerry: Sure! Amanda: I'll bring you tomorrow :-)"	"Amanda baked cookies and will bring Jerry some tomorrow."		
"Olivia: Who are you voting for in this election? Oliver: Liberals as always. Olivia: Me too!! Oliver: Great"	"Olivia and Olivier are voting for liberals in this election. "		
"Tim: Hi, what's up? Kim: Bad mood tbh, I was going to do lots of stuff but ended up procrastinating Tim: What did...	"Kim may try the pomodoro technique recommended by Tim to get more stuff done."		

Source: <https://huggingface.co/datasets/samsum>, <https://github.com/google-research/FLAN/blob/2c79a31/flan/v2/templates.py#L3285>

Sample FLAN-T5 prompt templates

```
"samsun": [  
  ("{dialogue}\n\nBriefly summarize that dialogue.", "{summary}"),  
  ("Here is a dialogue:\n{dialogue}\n\nWrite a short summary!",  
   "{summary}"),  
  ("Dialogue:\n{dialogue}\n\nWhat is a summary of this dialogue?",  
   "{summary}"),  
  ("{dialogue}\n\nWhat was that dialogue about, in two sentences or less?",  
   "{summary}"),  
  ("Here is a dialogue:\n{dialogue}\n\nWhat were they talking about?",  
   "{summary}"),  
  ("Dialogue:\n{dialogue}\n\nWhat were the main points in that "  
   "conversation?", "{summary}"),  
  ("Dialogue:\n{dialogue}\n\nWhat was going on in that conversation?",  
   "{summary}"),  
]
```


Improving FLAN-T5's summarization capabilities

Further fine-tune FLAN-T5 with a domain-specific instruction dataset (**dialogsum**)

Datasets: knkarthick/**dialogsum** like 13

Tasks: Summarization Text2Text Generation Text Generation Languages: English Multilinguality: monolingual Size Categories: none

Language Creators: expert-generated Annotations Creators: expert-generated Source Datasets: original License: mit

Dataset card Files and versions Community 4

Dataset Preview

Split

train (12.5k rows)

id (string)	dialogue (string)	summary (string)
"train_0"	"#Person1#: Hi, Mr. Smith. I'm Doctor Hawkins. Why are you here today? #Person2#: I found it would be a good..."	"Mr. Smith's getting a check-up, and Doctor Hawkins advises him to have one every year. Hawkins'll give some..."
"train_1"	"#Person1#: Hello Mrs. Parker, how have you been? #Person2#: Hello Dr. Peters. Just fine thank you. Ricky..."	"Mrs Parker takes Ricky for his vaccines. Dr. Peters checks the record and then gives Ricky a vaccine."
"train_2"	"#Person1#: Excuse me, did you see a set of keys? #Person2#: What kind of keys? #Person1#: Five keys and a small foot ornament. #Person2#: What a shame! I didn't see them. #Person1#: Well, can you help me look for it? That's my first time here. #Person2#: Sure. It's my pleasure. I'd like to help you look for the missing keys. #Person1#: It's very kind of you. #Person2#: It's not a big deal. Hey, I found them. #Person1#: Oh, thank God! I don't know how to thank you, guys. #Person2#: You're welcome."	"#Person1#'s looking for a set of keys and asks for #Person2#'s help to find them."

Example support-dialog summarization

Prompt (created from template)

Summarize the following conversation.

Tommy: Hello. My name is Tommy Sandals, I have a reservation.

Mike: May I see some identification, sir, please?

Tommy: Sure. Here you go.

Mike: Thank you so much. Have you got a credit card, Mr. Sandals?

Tommy: I sure do.

Mike: Thank you, sir. You'll be in room 507, nonsmoking, queen bed.

Tommy: That's great, thank you!

Mike: Enjoy your stay!

Source: <https://huggingface.co/datasets/knkarthick/dialogsum/viewer/knkarthick--dialogsum/>

Summary before fine-tuning FLAN-T5 with our dataset

Prompt (created from template)

Summarize the following conversation.


Tommy: Hello. My name is Tommy Sandals, I have a reservation.

Mike: May I see some
...
...
...

Tommy: That's great, thank you!

Mike: Enjoy your stay!

Model



Completion (Summary)

Tommy Sandals has a reservation for a room at the Venetian Hotel in Las Vegas.

Adequate completion, but does not match human baseline.

Human baseline summary:

Tommy Sandals has got a reservation. Mike asks for his identification and credit card and helps his check-in.

Summary **after** fine-tuning FLAN-T5 with our dataset

Prompt (created from template)

Summarize the following conversation.

Tommy: Hello. My name is Tommy Sandals, I have a reservation.

Mike: May I see some

...


...

...

Tommy: That's great, thank you!

Mike: Enjoy your stay!

Model



Completion (Summary)

Tommy Sandals has a reservation and checks in showing his ID and credit card. Mike helps him to check in and approves his reservation.

*Better summary,
more-closely matches
human baseline.*

Fine-tuning with your own data

